

Rohitash Kator

Software Engineer – Backend, Cloud & Full-Stack

☎ +91-9785787351 | ✉ rohitashkmt@gmail.com | 🌐 /itsrkator | 📱 /itsrkator | 📍 Jaipur, India

Professional Summary

Backend Engineer with 4 years of experience building and testing microservices, APIs, and AI/LLM systems in influencer marketing, FinTech, network intelligence, and document processing, using JavaScript/TypeScript (Node.js, NestJS), Python, and Go. Hands-on with agentic AI and LLM infrastructure — LangGraph multi-agent systems, Model Context Protocol (MCP) servers, and multi-provider LLM integrations (AWS Bedrock, Azure AI, OpenAI). Built an LLM-powered document extraction pipeline (AWS Textract & OpenAI) that hit 95%+ accuracy and cut manual data entry by over 90%. Google Cloud Certified (Associate Cloud Engineer – ACE, Professional Data Engineer – PDE) and GitHub Copilot Certified Professional. Comfortable owning a service end to end and turning a codebase with little or no test coverage into something reliable and well tested.

Education

Bachelor of Technology (B.Tech), Computer Science and Engineering
National Institute of Technology (NIT), Andhra Pradesh | Tadepalligudem, India

Aug 2018 – Jun 2022
Full-Time

Professional Experience

Software Engineer
Talentica Software | Remote, Pune

Nov 2024 – Present
Full-Time

Client: Mist Systems (Juniper Networks / HPE) – AI, Agentic Systems & Multi-Codebase Unit Testing

Python, pytest, Go, LangGraph, FastMCP, AWS Bedrock, Azure AI, OpenAI, NATS, Kafka, Spark, CircleCI, Cursor, GitHub Copilot

Extended a unit-testing engagement into full ownership of test coverage across the AI and infrastructure portfolio – the LLM/agentic AI backend, an AI-powered network analytics service, an automated test-orchestration framework, a Spark-based knowledge graph pipeline, and production Go microservices – using Cursor and GitHub Copilot for scaffolding with rigorous manual review for edge cases and contract alignment.

- Implemented behavior-driven unit tests for a LangGraph-based agentic AI service, covering multi-agent graphs (troubleshooting, text-to-SQL, entity extraction, action execution), multi-provider LLM integrations (AWS Bedrock, Azure AI, OpenAI), and a FastMCP Model Context Protocol server.
- Implemented unit tests for an AI-powered network analytics and troubleshooting service spanning anomaly detection, graph-topology reasoning, event correlation, conversational QA, model serving, and caching.
- Wrote Go unit tests for the assistant's log-parsing and event-routing service, and configured a unified Go/Python coverage-reporting pipeline on CircleCI with AWS S3 artifact upload.
- Built pytest coverage for an automated network test-orchestration framework, including AP authentication, gateway LAN, FIPS MTR, test-trigger scheduling, and cloud-storage reporting.
- Implemented unit tests for a Spark-based knowledge graph and semantic-search pipeline covering the web crawler, data indexing, and semantic-search validation.
- Delivered Go unit-test suites for production infrastructure services – webhook-to-Kafka ingestion, a Kafka/Redis firmware-upgrade orchestrator, a repository proxy, and a zone-management service – asserting routing logic and boundary conditions rather than raw status codes.

Client: Triller – Julius Works Influencer Marketing & Campaign Platform

TypeScript, NestJS, GraphQL, Angular, MySQL, DynamoDB, Redis, RabbitMQ, Elasticsearch, Microservices, AWS ECS, Docker

Contributed as a full-stack engineer across multiple microservices of Julius Works, an enterprise influencer marketing SaaS platform used by global brands and agencies for influencer discovery, campaign management, social data ingestion across nine social platforms, analytics, and influencer payouts.

- Architected a GraphQL BFF reporting API (NestJS, Apollo) aggregating campaign data from MySQL, DynamoDB, and Redis across two downstream microservices, serving content summaries, influencer metrics, and Earned Media Value (EMV) reporting.
- Diagnosed and fixed campaign-lifecycle tracking logic that was silently dropping active trackings through status transitions, eliminating payout-calculation errors across the reporting service.
- Built an end-to-end influencer payout system integrating the Dots payments API for ACH onboarding and balance management, with Redlock-based distributed scheduling and RabbitMQ payout-state sync with retry and dead-letter queues.
- Implemented a dedicated Instagram-ingestion microservice with a four-mode fallback strategy (BrightData/Luminati, first-party data, Phyllo API) and added the migration analysis adopted as the platform's roadmap for retiring legacy scrapers.
- Shipped campaign-reporting and influencer-payout Angular SPAs powered by various typed GraphQL operations, plus a Bitly/Shopify-integrated core REST API with Elasticsearch-based influencer discovery.
- Maintained GitHub Actions CI/CD pipelines for Dockerized microservices and AWS ECS deployments, and strengthened rollout safety through feature flags and expanded unit/integration coverage.

Software Engineer

Eternalight Infotech | Remote, Pune

Feb 2024 – Nov 2024

Full-Time

Client: FamO – FinTech Family Investment Tracking Platform

TypeScript, Node.js, NestJS, TypeORM, PostgreSQL, React.js, Firebase, AWS (S3, X-Ray, CloudWatch), GrowthBook

Owned backend development across multiple modules of FamO, a multi-tenant family office wealth management platform for tracking investments across equities, mutual funds, real estate, and private deals. Delivered scalable, well-tested features optimized for high traffic and reliability.

- Designed a multi-step family onboarding engine and the CAS/CAM statement-processing pipeline – OTP-secured document access, scheduled scraping jobs, and Tally accounting-software sync – automating a previously manual back-office workflow.
- Built account and portfolio APIs supporting live price updates, family-level net-worth aggregation, and discrepancy detection between reported and actual broker holdings.
- Authored 100+ TypeORM migrations and seed scripts supporting iterative, zero-downtime schema evolution in production.
- Built a task-management and threaded-discussions system from schema design through delivery, unifying tasks, messages, and discussions into one data model with Firebase push notifications across all three surfaces.
- Integrated AWS X-Ray tracing and CloudWatch structured logging across auth, broker, and notification services, and diagnosed a multi-connection race condition causing intermittent production failures.
- Rolled out GrowthBook feature flagging for incremental releases and migrated hardcoded credentials to environment-based configuration across dev, staging, and production.

Client: FamO – Contract Note Parser (AI Document Processing)

Python, Flask, SQLAlchemy, PostgreSQL, AWS Textract, AWS X-Ray, AWS S3, OpenAI

Sole backend engineer on a standalone microservice extracting structured trade data from broker-issued contract note PDFs, replacing fully manual back-office data entry.

- Designed an LLM-powered extraction pipeline combining AWS Textract OCR with the OpenAI Assistants API to pull structured transaction data – stock, quantity, price, brokerage, taxes – from contract notes at 95%+ accuracy across multiple broker formats.
- Implemented a page-chunking strategy that split multi-page notes into smaller segments before AI extraction, cutting processing time and token cost while improving reliability on large documents.
- Authored and iteratively tuned broker-specific extraction prompts, correcting real-world discrepancies in brokerage-deduction and pay-in/pay-out calculations.
- Built pre-persistence validation and a PostgreSQL schema (SQLAlchemy, Flask-Migrate) that eliminated the majority of manual data-entry review; resolved a concurrency race condition and added AWS X-Ray tracing with S3-based document storage.

Cloud Engineer

SpringML Inc. (Acquired by Egen.ai) | Remote, Hyderabad

Jul 2022 – Nov 2023

Full-Time

- Designed and built **MGIC DocAI**, a full-stack AI-powered invoice-parsing system (Angular, Flask, Vertex AI Document AI, PostgreSQL) that automated structured extraction – vendor details, line items, totals – across heterogeneous, non-templated invoice layouts.
- Built a human-in-the-loop correction workflow with side-by-side document-vs-extracted-data review, feeding reviewer edits back into the pipeline as ground truth and reducing manual data entry.
- Built a full-stack presales demo for a state DOT crash-reporting platform (Angular, Flask, GCP App Engine) featuring a configurable rules engine, multi-agency data-pipeline ingestion, and automated PDF compliance-letter generation; containerized and deployed it to GCP App Engine with autoscaling (1–20 instances) for live client demos.
- Managed GCP IAM configurations enforcing least-privilege access across multiple client projects, and built Looker/Looker Studio dashboards delivering real-time operational insights to stakeholders.

Internship Experience

2021 – 2022

- **Software Engineer Intern, iLink Digital** (Jan – Mar 2022): Developed RESTful APIs (Python, Flask, Django) for a data-driven analytics platform and built real-time Grafana dashboards integrated with PostgreSQL and Elasticsearch.
- **Software Developer Intern, Jivrus Technologies** (May – Nov 2021): Built an online assessment proctoring feature with camera tracking, screen monitoring, and secure media storage, integrating Google Drive and Maps APIs.

Technical Skills

Languages: JavaScript, TypeScript, Python, Go, Java

Backend: Node.js, Express.js, NestJS, Fastify, Flask, Gin, REST APIs, GraphQL, WebHooks, WebSockets, Microservices

Databases & ORMs: PostgreSQL, MySQL, MongoDB, DynamoDB, Redis, Elasticsearch, TypeORM, Prisma, SQLAlchemy

Cloud & DevOps: AWS, GCP, Docker, Kubernetes, GitHub Actions, CircleCI

Messaging & Data: Kafka, RabbitMQ, NATS, Socket.IO, Apache Spark

Frontend: React.js, Angular, HTML5, CSS3, Bootstrap

Testing & Tools: Jest, Supertest, pytest, unittest, Git, JIRA, ClickUp, Postman, DBever, GitHub Copilot, Cursor, Claude

Key Projects

Multi-Agent Cloud Support System | *Python, FastAPI, LangGraph, LangChain, React, PostgreSQL, pgvector, Pinecone, LangSmith, AWS*

- Designed a production-grade multi-agent AI platform using a LangGraph StateGraph orchestrating nine specialized agents (intent classification, RAG retrieval/reranking, reasoning, synthesis, guardrails, action handling); built a full RAG pipeline with dual vector storage (pgvector + Pinecone), semantic reranking, and RAGAS evaluation, plus a React frontend with real-time WebSocket streaming and agent-graph visualization.

Payment Processing System | *NestJS, TypeORM, PostgreSQL, Authorize.Net SDK, Swagger, Jest*

- Built a NestJS payment gateway integrating the Authorize.Net Sandbox API across the full payment lifecycle (purchase, authorize, capture, refund, cancel) with a transaction audit trail and interactive Swagger docs, backed by 225+ unit tests and 21 E2E tests.

SkyHigh Airlines Digital Check-In System | *NestJS, PostgreSQL, TypeORM, Kubernetes, Docker*

- Developed a concurrent seat-reservation system using pessimistic row locking to prevent double-booking and cron-based seat-hold expiry, with a Kubernetes deployment manifest targeting sub-second P95 seat-map response time.

Real-Time Chat Application | *NestJS, Angular, Socket.IO, PostgreSQL, TypeORM, Google OAuth*

- Engineered a full-stack chat platform with a JWT-guarded Socket.IO gateway, multi-tab presence tracking, dual-token auth with Google OAuth 2.0, and a Jest suite enforcing 100% branch coverage.

Employee Management REST API | *Go, Gin, MongoDB, Clean/Hexagonal Architecture*

- Built a clean/hexagonal-architecture REST API in Go with strict domain-usecase-adapter layering, structured JSON logging (slog), request-ID middleware, and graceful shutdown.

Certifications

- GitHub Copilot Certified Professional
- Google Cloud Certified – Professional Data Engineer (PDE)
- Google Cloud Certified – Associate Cloud Engineer (ACE)